



Sir Syed University
of Engineering & Technology

Artificial Intelligence Lab

Project Report

Supervised by
Miss Fozia Khan

Table of Content

About Project.....4

 Project Title4

 Project Members4

Abstract..... Error! Bookmark not defined.

Introduction Error! Bookmark not defined.

 Background / Motivation Error! Bookmark not defined.

 Problem Statement Error! Bookmark not defined.

 Objectives..... Error! Bookmark not defined.

 Scope of the Project Error! Bookmark not defined.

 Significance..... Error! Bookmark not defined.

Literature Review..... Error! Bookmark not defined.

 Existing Methods & Systems Error! Bookmark not defined.

 Gaps in Current Solutions..... Error! Bookmark not defined.

 Comparative analysis..... Error! Bookmark not defined.

Methodology / Proposed System..... Error! Bookmark not defined.

 Algorithm / Model..... Error! Bookmark not defined.

 Dataset Description Error! Bookmark not defined.

 Data Preprocessing Steps Error! Bookmark not defined.

 Model Architecture Error! Bookmark not defined.

 Tools and Technologies..... Error! Bookmark not defined.

Implementation Error! Bookmark not defined.

 Workflow Error! Bookmark not defined.

 Code structure / modules Error! Bookmark not defined.

 Key functions or pseudocode Error! Bookmark not defined.

 Complete GUI Overview10

 Screenshots or UI design11

Results and Evaluation..... 12-20

 Performance Metrics of Brain tumor model Error! Bookmark not defined.

 Performance Metrics of Lung Cancer model.....16

 Comparative Analysis: Brain Tumor Detection vs Lung Cancer Detection Models19

12. Conclusion.....21

 Summary21

 Achievements.....21

 Real-world Application Potential.....21

Future Work21

 Improvements21

Scalability ideas21

Integration possibilities21

Project Code.....22

References.....22

About Project

This project is a web-based AI-powered medical image analysis system that detects brain tumors and lung diseases from medical images (MRI and chest X-rays respectively). It leverages deep learning models trained using VGG16 and custom CNN architectures and integrates Gemini-powered agents to provide condition explanations, treatment plans, and doctor recommendations.

Project Title

Medical Image Analysis

Project Members

S No.	Roll No.	Sec	Name	Contact No.	Email
1	2022F-BCS-152	A	Muhammad Faisal	03013301586	muhammadfaisal159@gmail.com
2	2022F-BCS-162	A	Saad Ali	03182007567	m.saadaliarain@gmail.com
3	2022F-BCS-316	A	Muzamil Mehmood	03323092223	konnnectmuzamil@gmail.com



2022F-CS-152



2022F-CS-162



2022F-CS-316

Abstract

The growing burden of brain tumors and lung diseases, particularly pneumonia, poses a significant challenge to the global healthcare system. Early and accurate detection of these conditions plays a vital role in reducing mortality rates and improving patient outcomes. However, manual diagnosis through radiological imaging is time-consuming, prone to human error, and highly dependent on specialist availability. In response to these limitations, this project introduces a comprehensive deep learning-based medical image analysis system capable of detecting brain tumors and lung diseases from MRI and chest X-ray images, respectively.

This system integrates advanced Convolutional Neural Network (CNN) models with a user-friendly web interface to assist patients and healthcare providers in preliminary diagnosis. The brain tumor detection component leverages transfer learning using the pre-trained VGG16 architecture, while the lung disease detection model is built using a custom-designed CNN. Both models have been trained on publicly available datasets, achieving high classification performance and generalizability.

To further enhance the application's utility, the system incorporates AI-powered agents using the Gemini API to provide additional patient guidance. These agents can generate detailed condition explanations, propose customized treatment plans, and recommend local doctors based on the user's city input. This intelligent interaction makes the system not only diagnostic but also educational and supportive in nature.

The complete application is deployed using Flask and is fully accessible via a web browser. The backend models are hosted on Hugging Face, and the system is containerized using Docker for scalability and ease of deployment. A clean and responsive front-end interface allows users to upload their scans, select the detection type, and receive detailed results and guidance with high confidence scores.

The significance of this project lies in its ability to democratize access to preliminary diagnostic support, especially in regions with limited medical infrastructure. It showcases the power of integrating deep learning with real-time user interaction and intelligent agents in building a practical, cost-effective, and impactful medical tool. Through automation and intelligent reporting, the project has the potential to accelerate the diagnostic pipeline, support clinical decision-making, and pave the way for AI-assisted medical solutions in under-resourced areas.

Introduction

Background / Motivation

Accurate and timely diagnosis of life-threatening conditions like brain tumors and pneumonia remains a major challenge, especially in resource-limited settings. According to the World Health Organization (WHO), both brain tumors and lower respiratory infections rank among the top global causes of death. Diagnosing these conditions requires expert interpretation of medical images (e.g., MRIs and chest X-rays), which is time-consuming and often unavailable in remote areas due to a shortage of trained radiologists.

Artificial Intelligence (AI), particularly deep learning with Convolutional Neural Networks (CNNs), offers an efficient alternative by enabling automated, high-accuracy medical image classification. Transfer learning techniques further enhance performance, making expert-level detection accessible across broader regions. Additionally, intelligent agents (e.g., Gemini API) can improve patient interaction by providing explanations, treatment suggestions, and local doctor recommendations.

Problem Statement

Manual diagnosis of brain tumors and pneumonia is time-consuming, expertise-dependent, and error-prone. Lack of access to radiological services—especially in rural and under-resourced regions—leads to delays and poorer outcomes. There is a need for a reliable, automated, and scalable diagnostic system.

Objectives

This project aims to address these challenges through the development of an AI-based medical diagnostic system with the following core objectives:

- **Automated Detection:** Implement deep learning models (CNNs) to classify MRI scans for brain tumors and chest X-rays for lung infections (pneumonia).
- **Explainability & Support:** Integrate intelligent agents using the Gemini API to explain the diagnosed condition in layman's terms, suggest treatment plans, and recommend doctors in the user's city.
- **Web Accessibility:** Build an intuitive and responsive web interface where users can upload medical scans and receive detailed, understandable analysis in real-time.

Scope of the Project

This project includes:

- **Brain Tumor Detection:** Classification of brain MRI images into four categories: glioma, meningioma, pituitary tumor, and no tumor.
- **Lung Disease Detection:** Classification of chest X-ray images into pneumonia and normal categories.
- **AI Agents for Patient Guidance:** Automatic generation of markdown-formatted treatment plans, condition explanations, and doctor/hospital recommendations using location-based input.
- **Deployment:** Integration of the trained models into a Flask-based web application with models hosted on Hugging Face and deployed via Docker for universal accessibility.

Significance

The proposed system has wide-reaching implications:

- **For Patients:** Offers an initial assessment tool that can reduce anxiety, provide educational support, and motivate timely medical consultation.
- **For Healthcare Providers:** Acts as a decision support system, offering a second opinion and improving diagnostic confidence.
- **For Public Health:** Facilitates early detection in areas with limited access to medical expertise, thereby contributing to better health outcomes and reducing the burden on tertiary healthcare facilities.

This project makes medical AI not only technically effective but practically usable and socially impactful.

Literature Review

Existing Methods & Systems

1. Deep Learning for Brain Tumor Detection in MRI Images (2017, Pereira et al.)
This study proposed a CNN-based segmentation model to identify and classify brain tumors in MRI scans. The network achieved high accuracy in detecting tumor boundaries and types, laying the groundwork for automated tumor analysis. However, the system was not integrated into a user-facing application.
2. ChestX-ray8: Hospital-scale Chest X-ray Database and Benchmarks (2017, Wang et al.)
This influential work introduced a large-scale dataset and benchmarked deep learning models for multi-label chest pathology classification. Though powerful, the model focused on multiple diseases, not pneumonia specifically, and lacked explainability features for patient communication.
3. Transfer Learning with VGG16 for Pneumonia Detection (2019, Kermamy et al.)
Kermamy et al. applied VGG16 via transfer learning to detect pneumonia from chest X-rays. Their model achieved high accuracy (>92%) but was not part of an end-to-end system including treatment suggestions or web deployment.
4. Brain Tumor Classification Using CNN on MRI Images (2020, Swati et al.)
This paper employed a CNN trained from scratch to classify MRI brain scans into tumor types. While the approach performed well, it required a large dataset and extensive training and lacked real-time deployment in clinical settings.
5. Explainable AI for Medical Imaging: A Survey (2022, Holzinger et al.)
This survey reviewed the need for explainability in medical imaging AI. It highlighted that most models act as "black boxes" and fail to provide insights or treatment options for non-expert users, emphasizing the demand for integrated AI agents.

Gaps in Current Solutions

Feature	Observation
Explainability	Most models lack explanation for the prediction
Treatment Recommendations	Absent in nearly all academic approaches
Doctor/Location Integration	No personalization based on user’s city or context
Web Interface	Few research models offer user-accessible online tools
Unified Multi-Disease System	Systems often focus on one disease (brain or lung only)

Comparative Analysis

Feature	Research Papers	Our Project
Brain Tumor Detection	✓ (1, 4)	✓
Pneumonia Detection	✓ (2, 3)	✓
Transfer Learning (VGG16)	✓ (3)	✓
Explanation for Users	✗	✓ (Gemini API)
Doctor Recommendations	✗	✓ (City-based using Gemini agent)
Web-Based Deployment	✗	✓ (Flask + Docker + Hugging Face)

Methodology / Proposed System

Algorithm / Model

- **Brain Tumor Detection:** Utilizes transfer learning with the VGG16 architecture pre-trained on ImageNet. The base layers are frozen to retain learned features, and a custom classifier head with Global Average Pooling and dense layers is added to classify MRI images into glioma, meningioma, pituitary tumor, or no tumor.
- **Lung Disease Detection:** A custom Convolutional Neural Network (CNN) is designed with multiple convolutional and pooling layers, tailored to extract relevant features from chest X-rays and classify them into normal or pneumonia.
- **Gemini Agents :** The Gemini API is integrated to generate AI-based responses such as a plain-language explanation of the diagnosis, treatment steps, and location-specific doctor recommendations, making the system interactive and patient-friendly.

Dataset Description

- **Brain Tumor Dataset:** Consists of MRI scans organized into Training and Testing folders labeled as glioma, meningioma, pituitary tumor, and no tumor.
- **Lung Disease Dataset:** Chest X-ray images sorted into train/val/test directories with two classes: normal and pneumonia.

Data Preprocessing Steps

- **Resizing:**
 - Brain MRI: resized to 128x128 pixels
 - Chest X-rays: resized to 150x150 pixels
- **Data Augmentation:**
 - Applied to training data for generalization
 - Includes rotation ($\pm 20^\circ$), zoom, width/height shift, and brightness adjustment
- **Normalization:**
 - Pixel values scaled to [0, 1] using /255 to standardize input across models

Model Architecture

- **Brain Tumor Model (VGG16-Based):**
 - Frozen VGG16 base
 - Global Average Pooling layer
 - Dense layers with ReLU activation and Dropout
 - Final output layer with Softmax activation for 4-class classification
- **Lung Disease Model (Custom CNN):**
 - Three convolutional layers with ReLU and MaxPooling
 - Flatten layer followed by Dense(128) and Dropout(0.3)
 - Output layer with Sigmoid activation for binary classification

Tools and Technologies

- **Programming Language:** Python (3.10)
- **Deep Learning Frameworks:** TensorFlow 2.x, Keras
- **Web Development:** Flask (for backend), HTML/CSS + Bootstrap (for frontend)
- **AI Agents :** Gemini API & OpenAI SDK
- **Model Hosting:** Hugging Face Hub (for deploying and accessing trained models)
- **Environment:** Google Colab (model training), Docker (deployment), GitHub (code repo)

Implementation

Workflow

The project follows a structured, user-centric pipeline from image input to diagnostic output, as outlined below:

1. User Interaction:
 - The user visits the web interface built with Flask and Bootstrap.
 - They select the **detection type** (brain tumor or lung disease), enter their city, and upload an MRI or chest X-ray image.
2. Image Processing and Prediction:
 - Depending on the selected detection type:
 - Brain MRI images are fed into the VGG16-based model trained to detect glioma, meningioma, pituitary tumors, or no tumor.
 - Chest X-rays are processed through the custom CNN model to classify as pneumonia or normal.
 - The image is resized and normalized before being passed to the model.
3. AI Agent Interaction:
 - Once a prediction is made, Gemini API functions are triggered to:
 - Explain the detected condition in simple, understandable language.
 - Suggest a treatment plan relevant to the diagnosis.
 - Recommend local doctors or hospitals based on the user's city.
4. Result Display:
 - The final output includes:
 - Prediction label
 - Confidence percentage
 - Uploaded image preview
 - Condition explanation
 - Suggested treatment plan
 - Doctor/hospital recommendations
 - All information is rendered dynamically using Jinja2 templates in index.html.

Code Structure / Modules

File	Purpose
main.py	Flask backend containing routes, model loading, and image prediction logic
agent_helper.py	Contains Gemini agent functions for NLP responses like treatment plans
index.html	Frontend template with Bootstrap UI and Jinja2 rendering for dynamic output
Brain-Tumor-Detection.ipynb	Notebook to train the VGG16-based model for brain tumor classification
Lung-Cancer-Detection.ipynb	Notebook to build and train the CNN model for chest X-ray classification

Key Functions / Pseudocode

1. Prediction Functions:
 - predict_brain(image_path):
 - Loads and preprocesses the brain MRI image
 - Passes it to the VGG16-based model
 - Returns predicted class and confidence
 - predict_lung(image_path):
 - Loads and resizes the chest X-ray
 - Runs it through the custom CNN
 - Returns classification (pneumonia/normal) with confidence

2. Gemini-Powered Agents (agent_helper.py):
 - get_condition_explanation(condition):
 - Uses Gemini API to generate a patient-friendly explanation for the detected condition.
 - get_treatment_plan(condition):
 - Provides a markdown-formatted treatment strategy based on the diagnosis.
 - get_doctor_recommendation(condition, city):
 - Recommends medical experts or hospitals using location-based search logic and Gemini response generation.

Complete GUI Overview

The user interface of our Medical Image Analysis System is designed for ease of use, clarity, and clinical relevance. The workflow allows users to seamlessly upload patient imaging and receive immediate AI-driven diagnostic insights.

Interface Functionality:

1. Detection Type Selection

Users can choose between:

- Brain Tumor Detection
- Lung Cancer Detection

2. City Name Entry

The user enters their current city to enable the system to provide localized healthcare recommendations.

3. Image Upload

The interface allows users to upload medical images in .JPG or .PNG format. Depending on the selected detection type:

- MRI scans are used for Brain Tumor Detection.
- Chest X-rays are used for Lung Cancer Detection.

4. Action Button – "Detect & Analyze"

Once the image is uploaded, the user clicks the button to initiate analysis.

Screenshots or UI design


Medical Image Analysis

The screenshot displays a web interface for 'Medical Image Analysis'. It features a white form box on a light gray background. Inside the form, there are three main sections: 1. 'Select Detection Type:' with a dropdown menu currently showing 'Brain Tumor Detection'. 2. 'Enter Your City:' with a text input field containing 'Karachi'. 3. 'Upload Image (JPG, PNG):' with a file upload area showing a 'Choose File' button and the filename 'Te-pi_0010.jpg'. At the bottom of the form is a prominent blue button labeled 'Detect & Analyze'.

Prediction Result

Result: Tumor: pituitary

Confidence: 100.00%



Condition Explanation

What a Pituitary Tumor Is

A pituitary tumor is an abnormal growth of cells in the pituitary gland.

- The pituitary gland is a small, pea-sized gland located at the base of the brain.
- It's often called the "master gland" because it controls many important functions in the body by releasing hormones. These hormones affect:
 - Growth
 - Blood pressure
 - Reproduction
 - And other vital processes
- Pituitary tumors are generally benign (non-cancerous). This means they don't spread to other parts of the body.
- However, even though they're not cancerous, they can still cause problems. They can do this by:
 - Pressing on nearby structures in the brain
 - Causing the pituitary gland to make too much or too little of certain hormones

Common Symptoms

Symptoms of a pituitary tumor can vary depending on the size of the tumor and which hormones are affected. Some common symptoms include:

- **Headaches:** Persistent or severe headaches
- **Vision Problems:** Blurred vision, double vision, or loss of peripheral vision.
- **Hormone-Related Symptoms:** These vary widely based on the specific hormone imbalances. Examples include:
 - **In women:** Irregular menstrual cycles, infertility, breast milk production when not pregnant or breastfeeding.
 - **In men:** Erectile dysfunction, loss of body hair.
 - **In both men and women:** Weight gain, fatigue, muscle weakness, increased blood sugar, increased body hair, changes in mood, and increased size of hands and feet (acromegaly).
 - **Nausea and Vomiting:** Especially if the tumor is large.

Causes and Risk Factors

- **Cause:** The exact cause of pituitary tumors is usually unknown. Most pituitary tumors arise spontaneously (without a clear reason).
- **Genetics:** In rare cases, pituitary tumors can be associated with genetic conditions like Multiple Endocrine Neoplasia type 1 (MEN1).
- **Risk factors:** There are not many established risk factors. Family history of certain genetic conditions may increase the risk.

How Serious It Is

The seriousness of a pituitary tumor depends on several factors, including:

- **Size:** Larger tumors can cause more pressure on nearby structures.
- **Hormone Production:** Whether the tumor is causing the pituitary gland to produce too much or too little of certain hormones.
- **Overall Health:** A person's general health and other medical conditions can also play a role.

While pituitary tumors are usually benign, they can cause significant health problems if they are:

- Left untreated
- Interfere with vision
- Cause hormone imbalances that affect important bodily functions

Treatment Plan

What To Do First

- **Schedule an Appointment:** See an endocrinologist or neurosurgeon for a full evaluation.
- **Further Testing:** You may need blood tests, MRI or CT scans to understand the tumor's size and how it affects your hormones.

Who To Consult

- **Endocrinologist:** Specializes in hormone disorders and pituitary function.
- **Neurosurgeon:** A surgeon who can operate on the brain and pituitary gland.
- **Ophthalmologist:** To check your vision, as pituitary tumors can affect eyesight.

Possible Medications

- **Hormone-Blocking Drugs:** If the tumor is producing too much of a hormone (like prolactin), medication can help lower the levels.
- **Hormone Replacement:** If the tumor is affecting normal hormone production, you might need hormone supplements.

How To Recover Faster

- **Follow Medical Advice:** Stick to the treatment plan your doctors give you.
- **Healthy Lifestyle:** Eat well, exercise, and get enough sleep to support your overall health.
- **Stress Management:** Use relaxation techniques to lower stress, as stress can affect hormone levels.

When To Get Help

- **Severe Headaches:** A sudden, severe headache could indicate a problem.
- **Vision Changes:** Any new or worsening vision issues need immediate attention.
- **Unexplained Symptoms:** If you have new symptoms or your current symptoms get worse, talk to your doctor.

Prevention Tips

- **Regular Check-ups:** While you can't prevent a pituitary tumor, regular check-ups can help catch any issues early.
- **Healthy Habits:** Maintain a balanced diet and regular exercise to support your body's overall health.

Recommended Doctors & Hospitals

- **Dr. Nadeem Khan** – Neuroendocrinology – Clifton – Specializes in pituitary disorders.
- **Aga Khan University Hospital** – Neurosurgery – Karachi – Comprehensive pituitary tumor treatment.
- **Dr. Aisha Memon** – Endocrinology – DHA – Focuses on hormonal imbalances due to tumors.
- **Liaquat National Hospital** – Endocrinology – Karachi – Offers advanced neuroendocrine services.

Results and Evaluation

Performance Metrics of Brain tumor model

Classification Report:				
	precision	recall	f1-score	support
glioma	0.94	0.90	0.92	300
meningioma	0.97	0.78	0.86	306
notumor	0.95	1.00	0.97	405
pituitary	0.84	0.99	0.91	300
accuracy			0.92	1311
macro avg	0.93	0.92	0.92	1311
weighted avg	0.93	0.92	0.92	1311

Precision: How many of the predicted positives are actually correct.

Recall: How many of the actual positives were correctly predicted.

F1-score: Harmonic mean of precision and recall, a balanced measure.

Support: Number of actual occurrences of the class in the test data.

Model Accuracy

- **Accuracy = 92%**

It means **the model correctly classified 92%** of the total 1,311 test samples.


Averages

Metric	Precision	Recall	F1-Score
Macro avg	0.93	0.92	0.92
Weighted avg	0.93	0.92	0.92

- **Macro avg:** Average of each class metric (treats all classes equally).
- **Weighted avg:** Weighted average based on support (more representative for imbalanced datasets).

Summary

- The model performs **best** on the **notumor** class (F1-score = 0.97).
- It performs **least well** on **meningioma** (Recall = 0.78, F1 = 0.86), possibly due to overlapping features with other tumor types.
- Overall, the model has a **strong performance** with balanced metrics and **92% accuracy**.



- **Rows** represent the actual class labels.
- **Columns** represent the predicted labels by the model.

Correct Predictions (True Positives)

- **Glioma** correctly classified: **271**
- **Meningioma** correctly classified: **238**
- **Notumor** correctly classified: **404**
- **Pituitary** correctly classified: **297**

Misclassifications

- **Glioma** was misclassified as:
 - Meningioma: 6
 - Notumor: 5
 - Pituitary: 18
- **Meningioma** was misclassified as:
 - Glioma: 15
 - Notumor: 14
 - Pituitary: 39
- **Notumor** misclassified as Glioma: 1
- **Pituitary** misclassified as:
 - Meningioma: 2
 - Notumor: 1


Observations

- **Best performance:** Notumor (404 out of 405 correctly classified → **~99.75% accuracy**)
- **Most confusion:** Meningioma (only 238 out of 306 correctly classified → **~77.78% accuracy**)
 - Most frequently confused with **Pituitary** and **Glioma**
- **Overall Model Accuracy:**

$$\frac{271 + 238 + 404 + 297}{1311} = \frac{1210}{1311} \approx 0.9229 \approx \mathbf{92.29\%}$$

Conclusion

The confusion matrix shows that while the model is highly effective overall—especially in detecting "Notumor" cases—it could be further improved by focusing on better distinguishing **Meningioma** from **Pituitary** and **Glioma**.



ROC Curve Analysis

The **Receiver Operating Characteristic (ROC) Curve** is a graphical representation of a classifier's performance across different threshold levels. It plots the **True Positive Rate (Recall)** against the **False Positive Rate**, helping evaluate how well the model distinguishes between classes.

ROC Curve Summary

In multi-class classification (like this one), the ROC curve is plotted for each class against all other classes (one-vs-rest). Below are the AUC (Area Under the Curve) values for each class:

Class	AUC Score
Glioma	0.99
Meningioma	0.99
Notumor	1.00
Pituitary	1.00

AUC values range from 0 to 1, where **1.0 indicates perfect classification** and **0.5 indicates random guessing**.

Interpretation

- The curves for all four classes are **very close to the top-left corner**, indicating **high sensitivity** and **low false positive rates**.
- The **AUC of 1.00** for both **Notumor** and **Pituitary** shows that the model is nearly perfect at distinguishing these classes.
- **Glioma** and **Meningioma** also perform excellently, with **AUC = 0.99**.

Key Insights

- The **model demonstrates exceptional classification performance** for all tumor types.
- The high AUC values confirm the reliability and robustness of the model's predictions.
- ROC curves complement the confusion matrix and classification report by visually verifying the model's ability to separate classes effectively.

Performance Metrics of Lung cancer model

Classification Report:				
	precision	recall	f1-score	support
0	0.98	0.62	0.76	234
1	0.81	0.99	0.89	390
accuracy			0.85	624
macro avg	0.90	0.81	0.83	624
weighted avg	0.88	0.85	0.85	624

Classification Report for Lung Cancer Detection

The classification report provides a detailed performance analysis of the trained model in distinguishing between two classes:

- **Class 0:** Non-cancer (normal)
- **Class 1:** Lung cancer (positive)

Class-wise Metrics:

Class	Precision	Recall	F1-Score	Support
0 (Normal)	0.98	0.62	0.76	234
1 (Cancer)	0.81	0.99	0.89	390



- **Precision:** Out of all predicted samples for a class, how many were actually correct.
 - Class 0: Very high precision (0.98) – most predicted "normal" cases were indeed normal.
 - Class 1: Lower precision (0.81) – some "cancer" predictions were false positives.
- **Recall:** Out of all actual samples of a class, how many were correctly predicted.
 - Class 0: Lower recall (0.62) – the model missed some actual normal cases (false negatives).
 - Class 1: Very high recall (0.99) – almost all actual cancer cases were correctly detected.
- **F1-score:** Harmonic mean of precision and recall.
 - Class 0: 0.76
 - Class 1: 0.89
 - Indicates the model is more effective in detecting **cancer** than non-cancer cases.

Overall Metrics:

- **Accuracy:** 85% – the overall percentage of correct predictions.
- **Macro Average:** Averages metrics equally across both classes.
 - Precision: 0.90, Recall: 0.81, F1-score: 0.83
- **Weighted Average:** Averages metrics weighted by the number of instances in each class.
 - Precision: 0.88, Recall: 0.85, F1-score: 0.85

Interpretation:

- The model performs **very well in identifying cancer cases (class 1)**, with high recall and F1-score.
- However, it is **less effective in correctly identifying normal cases (class 0)**, indicating a higher false positive rate.
- This tradeoff might be acceptable in a medical setting, where **missing a cancer case (false negative)** is riskier than **flagging a normal case as cancer (false positive)**.



Model Performance Evaluation

The figures above show the **accuracy** and **loss** curves of the convolutional neural network (CNN) model used for lung cancer detection, plotted over 10 training epochs.

1. Accuracy Plot (Left)


- The **blue line** represents the **training accuracy**, which steadily increases from around **83%** to **95%**, indicating the model is learning effectively on the training data.
- The **orange line** represents the **validation accuracy**, which fluctuates but generally remains high, peaking at around **94%**.
- Although some **instability** is observed in the validation accuracy (e.g., sudden drops), it recovers quickly. This might be due to variability in the validation dataset or minor overfitting.

2. Loss Plot (Right)

- The **training loss** (blue line) shows a consistent **downward trend**, starting around **0.38** and ending below **0.13**, indicating that the model is minimizing errors on the training set.
- The **validation loss** (orange line) fluctuates significantly, with sharp spikes, especially around epochs 1, 5, and 7. These variations suggest some **overfitting**, where the model performs well on training data but inconsistently on unseen validation data.

Interpretation

- The **high training accuracy and low loss** confirm that the model has successfully learned from the training data.
- However, the **variability in validation metrics** implies that the model's generalization ability might be limited and can be improved by:
 - Introducing regularization techniques (e.g., Dropout or L2 Regularization)
 - Using more diverse or larger validation data
 - Applying early stopping to prevent overfitting



Confusion Matrix Analysis

The confusion matrix above provides a visual representation of the model's performance in classifying between two classes:

- **Normal (Non-cancer)**
- **Pneumonia (interpreted here as Lung Cancer Positive)**

	Predicted: Normal	Predicted: Pneumonia
Actual: Normal	146 (True Negative)	88 (False Positive)
Actual: Pneumonia	3 (False Negative)	387 (True Positive)

Key Insights:

- **True Positives (TP):** 387
→ Cancer cases correctly predicted as cancer.
- **True Negatives (TN):** 146
→ Normal cases correctly predicted as normal.
- **False Positives (FP):** 88
→ Normal cases incorrectly predicted as cancer (potentially alarming but less risky than FN in medical diagnosis).
- **False Negatives (FN):** 3
→ Cancer cases incorrectly predicted as normal (a serious error in a clinical setting).

Interpretation:

- The model demonstrates **strong ability in identifying pneumonia/lung cancer**, with **only 3 false negatives**, indicating high **recall** for the positive class.
- However, the model produced **88 false positives**, classifying many normal cases as pneumonia. This reflects a **trade-off** toward **sensitivity** over **specificity**, which is acceptable in medical applications where **missing a disease case is more critical than a false alarm**.

Performance Balance:

- Ideal in early detection systems where **cancer detection is prioritized**.

- **Further refinement** (e.g., threshold tuning or using ensemble models) may help reduce false positives and improve reliability.

Comparative Analysis: Brain Tumor Detection vs Lung Cancer Detection Models

This section presents a comprehensive comparison between two deep learning models developed for medical image classification tasks—**Brain Tumor Detection** using MRI scans and **Lung Cancer Detection** using chest X-ray images. The analysis covers their respective methodologies, model architectures, training strategies, evaluation metrics, and diagnostic effectiveness.

1. Objective and Dataset Overview

Aspect	Brain Tumor Detection	Lung Cancer Detection
Detection Goal	Classify MRI brain images into 4 categories (tumor types)	Binary classification of chest X-rays (cancer vs normal)
Input Modality	MRI brain scans	Chest X-rays
Classes	Glioma, Meningioma, Pituitary, Notumor	0: Normal, 1: Lung Cancer
Total Test Samples	1,311	624

2. Methodology and Model Architecture

Feature	Brain Tumor Model	Lung Cancer Model
Preprocessing Techniques	Image resizing, normalization, class balancing	Image normalization, augmentation
Architecture	Custom CNN + Transfer Learning (VGG16)	Convolutional Neural Network (CNN)
Transfer Learning	Yes, with VGG16 pretrained on ImageNet	No
Loss Function	Categorical Crossentropy	Binary Crossentropy
Optimizer	Adam	Adam
Epochs	10–15 (based on convergence)	10
Data Augmentation	Yes	Yes

3. Model Accuracy and Overall Performance

Metric	Brain Tumor Model	Lung Cancer Model
Accuracy	92%	85%
Best Class Accuracy	Notumor (≈99.75%)	Cancer (Recall = 0.99)
Challenging Class	Meningioma (≈78% accuracy)	Normal (Recall = 0.62)
Overfitting Signs	Minimal; stable metrics	Yes; validation loss fluctuates
AUC Scores	0.99–1.00 (Excellent)	Not calculated

4. Classification Metrics (Precision, Recall, F1-Score)

Brain Tumor Model (Multi-Class)

Metric Type	Precision	Recall	F1-Score
Macro Average	0.93	0.92	0.92
Weighted Avg	0.93	0.92	0.92

Lung Cancer Model (Binary)

Class	Precision	Recall	F1-Score	Support
0	0.98	0.62	0.76	234
1	0.81	0.99	0.89	390
Macro Avg	0.90	0.81	0.83	—
Weighted Avg	0.88	0.85	0.85	—

Interpretation:

- The brain tumor model shows **balanced precision and recall across multiple tumor types**, excelling particularly in Notumor and Pituitary classification.
- The lung cancer model demonstrates **high sensitivity (Recall = 0.99)** in detecting cancer, which is crucial in medical diagnosis, but lower specificity due to false positives in normal cases.

5. Confusion Matrix Insights

Model	Key Observations
Brain Tumor	- Best classification for Notumor (404/405 correct)- Meningioma is the most misclassified class, often confused with Pituitary and Glioma- Strong multi-class performance overall
Lung Cancer	- True Positives (Cancer): 387- False Negatives (Cancer missed): only 3 (very low)- False Positives (Normal misclassified): 88- Prioritizes sensitivity (cancer detection) over specificity

6. ROC Curve & AUC Performance (Brain Tumor Only)

- **AUC Scores:**
 - Glioma: 0.99
 - Meningioma: 0.99
 - Notumor: 1.00
 - Pituitary: 1.00

Interpretation:

- The AUC values confirm **outstanding discriminatory power** across all classes.
- High AUC (close to 1.0) supports the reliability of predictions.

7. Strengths and Limitations

Criteria	Brain Tumor Model	Lung Cancer Model
Strengths	- High multi-class accuracy- Strong ROC/AUC results- Effective use of VGG16	- Excellent cancer detection (high recall)- Simpler binary setup
Limitations	- Meningioma classification less accurate	- High false positives (normal predicted as cancer)- Some overfitting
Clinical Relevance	Well-suited for detecting specific tumor types	Ideal for early cancer screening (high sensitivity)

8. Comparative Summary

Aspect	Brain Tumor Detection	Lung Cancer Detection
Use Case	Differentiates multiple brain tumor types	Binary classification of cancer presence
Model Accuracy	92%	85%
Sensitivity (Recall)	High (especially for Notumor, Pituitary)	Extremely high for cancer (0.99)
False Positives	Low	Higher (88 cases)
False Negatives	Moderate (Meningioma misclassification)	Very low (only 3 cancer cases missed)
Transfer Learning	Yes (VGG16)	No
Model Stability	Stable training and validation curves	Some validation instability

9. Conclusion

Both models exhibit strong performance in their respective domains, with unique strengths tailored to the clinical importance of each task. The **Brain Tumor Detection model** achieves high classification accuracy across four tumor types, demonstrating robustness with the use of **transfer learning (VGG16)**. On the other hand, the **Lung Cancer Detection model** prioritizes sensitivity, ensuring nearly all cancer cases are detected, making it highly suitable for **early-stage cancer screening systems** where **minimizing false negatives** is paramount.

Conclusion

Summary

This project successfully demonstrates the development of a comprehensive AI-powered medical diagnostic system that integrates deep learning for image classification with generative AI agents for decision support. The platform classifies brain MRI and chest X-ray images with high accuracy using VGG16 and a custom CNN, respectively. Users interact with the system through a clean, responsive web interface built with Flask and Bootstrap. Additionally, the integration of the Gemini API allows the system to deliver natural language explanations, treatment guidance, and personalized doctor recommendations, transforming it into an intelligent diagnostic assistant.

Achievements

- Achieved high prediction accuracy in detecting brain tumors and lung infections.
- Combined vision (CNN models) and language (Gemini agents) to provide end-to-end analysis and medical advice.
- Successfully hosted deep learning models on Hugging Face and deployed the app using Docker, ensuring portability and accessibility.
- Demonstrated seamless interaction between a user-friendly frontend and a powerful backend with real-time response generation.
- Validated with multiple test cases using real MRI and X-ray images, showcasing practical utility.

Real-world Application Potential

- **Hospitals and Clinics:** As a diagnostic aid for radiologists to speed up decision-making.
- **Telemedicine Platforms:** Empower remote consultations by offering AI-generated assessments before doctor review.
- **Health Awareness Campaigns:** As a self-screening tool to increase public awareness about brain and lung diseases.
- **Primary Healthcare Units:** Especially beneficial in rural or under-resourced settings lacking specialist access.

Future Work

Improvements

- Extend the models to support **multi-label classification**, enabling the detection of multiple co-existing conditions.
- Improve accuracy through **larger datasets**, **transfer learning enhancements**, or **ensemble models**.
- Develop a **cross-platform mobile app** (Android/iOS) for broader reach and usability.

Scalability Ideas

- Add support for additional disease types: **COVID-19**, **Tuberculosis**, **Breast Cancer**, and **Skin Lesions**.
- Enable multilingual support using AI language translation models for better accessibility.
- Introduce **user accounts** and history tracking to allow follow-up comparisons and personal data logs.

Integration Possibilities

- Connect with **Electronic Health Record (EHR)** systems to store or fetch patient history for improved decision-making.
- Build a **WhatsApp or Telegram chatbot interface** to allow image uploads and receive analysis directly within chat platforms.
- Use **Google Maps API or local directories** to fetch real-time doctor and hospital availability based on the user's city.
- Integrate **voice interaction features** to make the system more accessible for visually impaired or elderly users.

Project Code and models

Available on GitHub and Hugging Face:

- <https://huggingface.co/faissal159/Brain-Tumor-Detection>
- <https://huggingface.co/faissal159/Lung-Cancer-Detection>
- <https://github.com/FaisalBasil/Medical-Image-Analysis>

References

1. **Pereira, S., Pinto, A., Alves, V., & Silva, C. A.** (2017). Brain tumor segmentation using convolutional neural networks in MRI images. *IEEE Transactions on Medical Imaging*, 35(5), 1240–1251.
<https://doi.org/10.1109/TMI.2016.2538465>
2. **Wang, X., Peng, Y., Lu, L., Lu, Z., Bagheri, M., & Summers, R. M.** (2017). ChestX-ray8: Hospital-scale chest X-ray database and benchmarks on weakly-supervised classification and localization of common thorax diseases. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2097–2106.
<https://doi.org/10.1109/CVPR.2017.369>
3. **Kermany, D. S., Zhang, K., & Goldbaum, M.** (2018). Identifying medical diagnoses and treatable diseases by image-based deep learning. *Cell*, 172(5), 1122–1131.e9.
<https://doi.org/10.1016/j.cell.2018.02.010>
4. **Swati, Z. N. K., Zhao, Q., Kabir, M., Ali, F., Jamil, A., & Lu, J.** (2020). Brain tumor classification for MR images using transfer learning and fine-tuning. *Computerized Medical Imaging and Graphics*, 75, 34–46.
<https://doi.org/10.1016/j.compmedimag.2019.101647>
5. **Holzinger, A., Langs, G., Denk, H., Zatloukal, K., & Müller, H.** (2019). Causability and explainability of artificial intelligence in medicine. *Wiley Interdisciplinary Reviews: Data Mining and Knowledge Discovery*, 9(4), e1312.
<https://doi.org/10.1002/widm.1312>
6. Kaggle Datasets: MRI and Chest X-ray